



Insertion in a B-Tree



$n = 2$

Insert: 62


Insertion in a B-Tree



$n = 2$


Insert: 62

Insertion in a B-Tree




Insert: 50

Insertion in a B-Tree




Insert: 50

Insertion in a B-Tree




Insert: 75

Insertion in a B-Tree




Insert: 75


Insertion




Insertion




Insertion




Insertion




Insertion




Insertion




Insertion



Insertion




Insertion



Insertion




Insertion




Insertion: Primitives

- Inserting into a leaf node
- Splitting a leaf node
- Splitting an internal node
- Splitting root node


Inserting into a Leaf Node



Inserting into a Leaf Node




Inserting into a Leaf Node




Splitting a Leaf Node


61




Splitting a Leaf Node




Splitting a Leaf Node




Splitting a Leaf Node




Splitting a Leaf Node



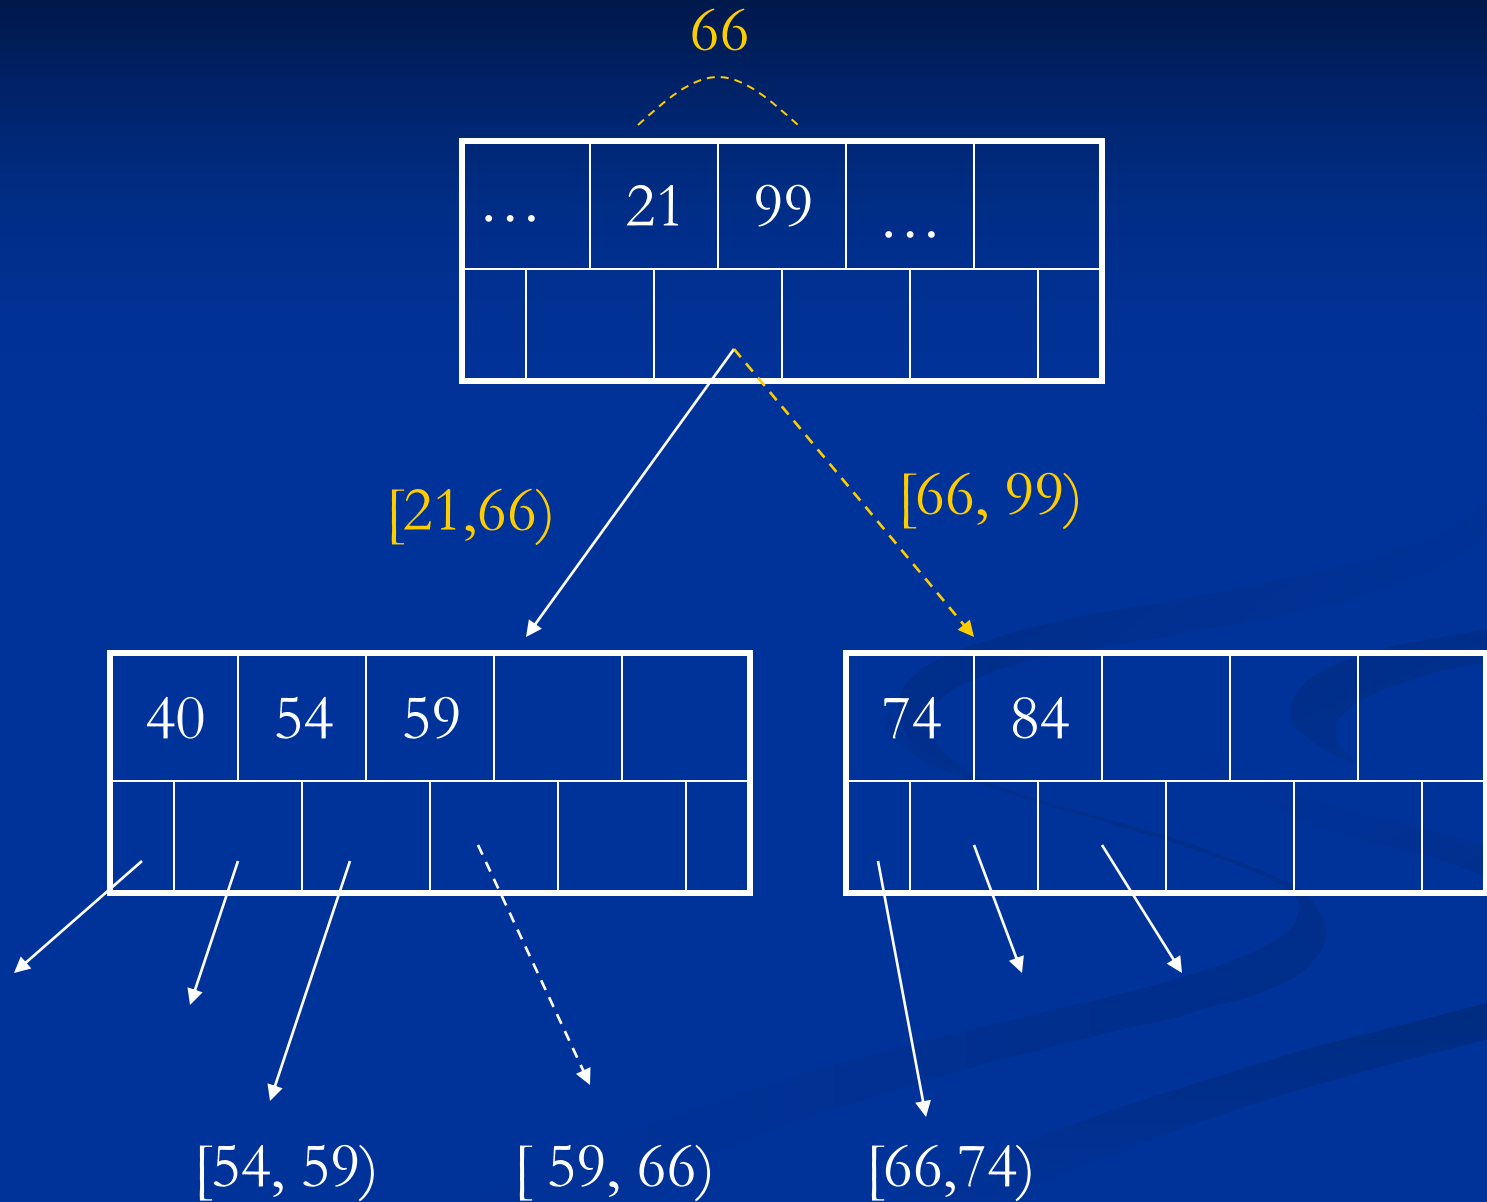
Splitting an Internal Node




Splitting an Internal Node




Splitting an Internal Node



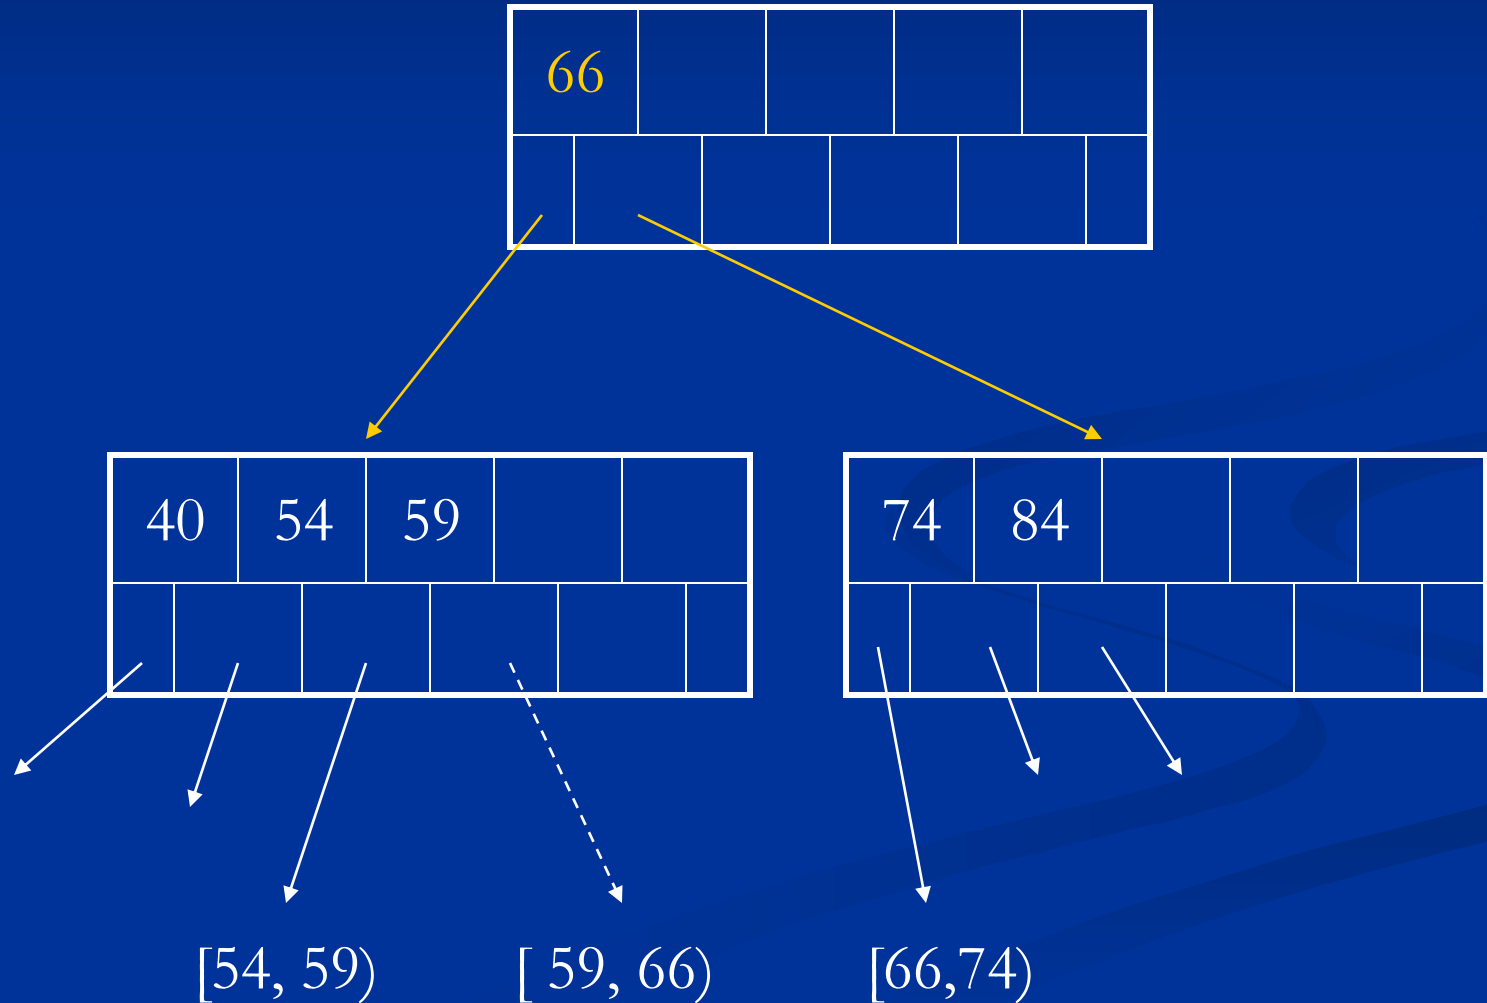
Splitting the Root




Splitting the Root



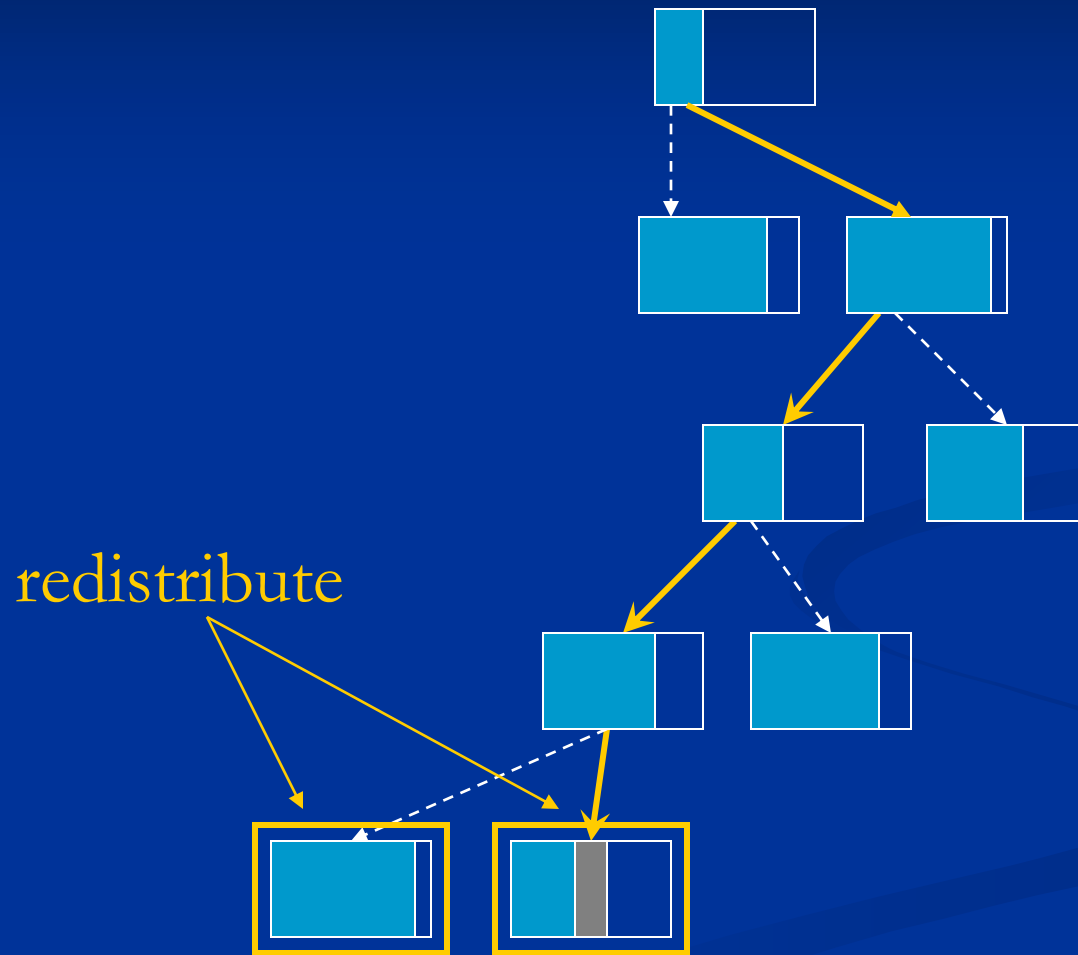
Splitting the Root




Deletion




Deletion




Deletion




Deletion - II




Deletion - II




Deletion - II



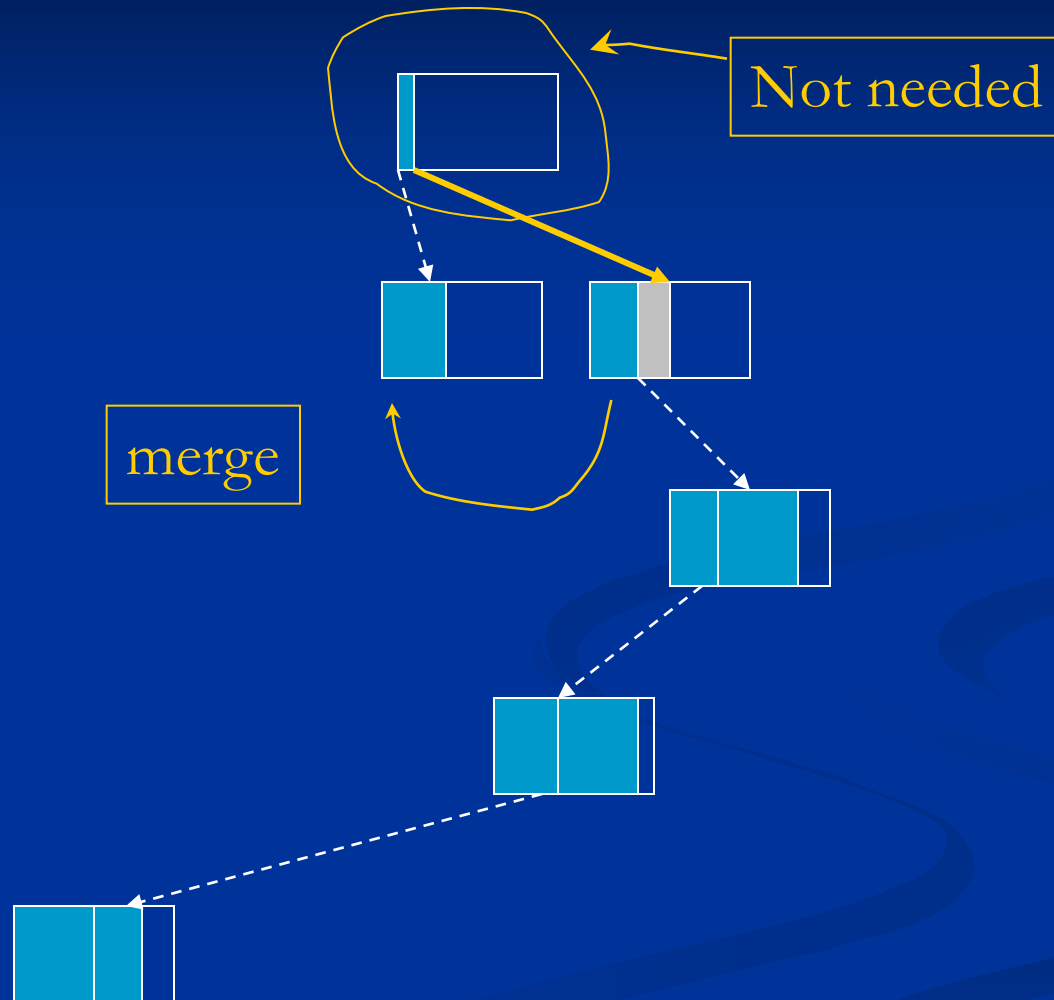
Deletion - II




Deletion - II



Deletion - II




Deletion - II




Deletion: Primitives

- Delete key from a leaf
- Redistribute keys between sibling leaves
- ⇒ ■ Merge a leaf into its sibling
- Redistribute keys between two sibling internal nodes
- ⇒ ■ Merge an internal node into its sibling


Merge Leaf into Sibling



Merge Leaf into Sibling



Merge Leaf into Sibling




Merge Leaf into Sibling

72


	...	85		

54	58	64	68	75

Merge Internal Node into Sibling



Merge Internal Node into Sibling



B-Tree Roadmap

- B-Tree

 - Recap

 - Insertion (recap)

 - Deletion

 - ➡ ■ Construction

 - Efficiency

- B-Tree variants

- Hash-based Indexes

Question

How does insertion-based construction perform?


B-Tree Construction

Sort




48 57 41 15 75 21 62 34 81 11 97 13

B-Tree Construction



Scan

B-Tree Construction



Scan

B-Tree Construction

Why is sort-based construction better than insertion-based one?

Cost of B-Tree Operations

- Height of B-Tree: H
- Assume no duplicates
- Question: what is the random I/O cost of:
 - Insertion:
 - Deletion:
 - Equality search:
 - Range Search:

Height of B-Tree

- Number of keys: N
- B-Tree parameter: n

$$\text{Height} \approx \log_n N = \frac{\log N}{\log n}$$

In practice: 2-3 levels

Question: How do you pick parameter n ?

1. Ignore inserts and deletes
2. Optimize for equality searches
3. Assume no duplicates

Roadmap

- B-Tree
- B-Tree variants
 - Sparse Index
 - Duplicate Keys
- Hash-based Indexes


Roadmap

- B-Tree
- B-Tree variants
- Hash-based Indexes
 - ➔ ■ Static Hash Table
 - Extensible Hash Table
 - Linear Hash Table


Hash-Based Indexes

- Adaptations of main memory hash tables
- Support equality searches
- No range searches

Indexing Problem (recap)




Main Memory Hash Table



Adapting to disk

- 1 Hash Bucket = 1 Block
 - All keys that hash to bucket stored in the block
 - Intuition: keys in a bucket usually accessed together
 - No need for linked lists of keys ...

Adapting to Disk




How do we handle this?

Adapting to disk

- 1 Hash Bucket = 1 Block
 - All keys that hash to bucket stored in the block
 - Intuition: keys in a bucket usually accessed together
 - No need for linked lists of keys ...
 - ... but need linked list of blocks (**overflow blocks**)

Adapting to Disk



Adapting to disk

- Bucket Id → Disk Address mapping
 - Contiguous blocks
 - Store mapping in main memory
 - Too large?
 - Dynamic ➔ Linear and Extensible hash tables

Beware of claims that assume 1 I/O
for hash tables and 3 I/Os for B-Tree!!