



Applications of Suffix Trees

Charles Yan
2008



1. Exact String Matching

$|P|=n, |T|=m$

- P and T are both known at the same time
 - Boyer-Moore, or Suffix trees. $O(n+m)$
- T is known and kept fixed. P varies.
 - Suffix trees, $O(m)$ in preprocess, $O(n+k)$ in searching
- P is known and kept fixed. T varies.
 - Boyer-Moore, $O(n)$ in preprocess, $O(m)$ in searching



2. Exact Set Matching

$|T|=m, P=\{p_1, p_2, \dots, p_i\}, \sum |p_i|=n$

- Aho-Corasick

$O(m+n+k)$

- Suffix trees.

$O(m)$ in building suffix tree

$O(n_i+k_i)$ in searching for p_i

$O(m+\sum n_i+ \sum k_i)$ for all P , i.e. $O(m+n+k)$



3. Substring problem for a set of texts

Motivation 1:

T is a DNA database containing millions of DNA sequences that have been previously sequenced.

Given a new DNA sequence, to determine whether it has been previously sequenced.

- (1) Concatenate all T together, then use Boyer-Moore $O(m+n+k)$ for searching each P, **m is huge!**
- (2) Build a suffix tree for each T_i
 $O(m)$ for total preprocessing, but $O(i * n+k)$ for searching each P, **i is in the order of 10^6 !**



Substring problem for a set of texts

Motivation 2:

To identify the remains of military personnel

For each soldier, a set of DNA sequences ($T = \{T_1, T_2, \dots, T_i\}$) is kept when he/she joins the army. (The whole genome sequence is very difficult to obtain for technical reasons.)

A DNA sequence (P) is extracted from the remains of personnel that have been killed.

To determine whether the remains belong to soldier A , we just need to see whether P matches any sequence in the T of A .



3. Substring problem for a set of texts

Given $T = \{T_1, T_2, \dots, T_i\}$, $\sum |T_i| = m$, $|P| = n$, set T is fixed, P varies. $O(m)$ preprocessing time is allowed. For each coming P , to find all occurrences of P in all T with $O(n+k)$ time

For each given P , this is the **reverse of exact set matching problem**.

(1) Concatenate all T together, then use Boyer-Moore

$O(m+n+k)$ for searching each P

(2) Build a suffix tree for each T_i ,

$O(m)$ for total preprocessing, but $O(i * n+k)$ for searching each P

(3) Build a suffix tree (**generalized suffix tree**) for the set T ,

the searching will take $O(n+k)$ time


but how to build the a generalized suffix tree in $O(m)$?



Generalized Suffix Trees

How to build the generalized suffix tree for a set $T = \{T_1, T_2, \dots, T_i\}$ in $O(m)$?

- (1) Append a marker to the end of each string and concatenated them together to build a new string S .
- (2) Build a suffix tree for S .
- (3) But, suffixes span multiple T_i ,





Generalized Suffix Trees


Minor subtleties

- (1) Each edge is associated with three indices (i, p, q) , where i indicates that the substring come from T_i . p and q are the begin and end positions.
- (2) Suffixes from two texts may be identical. Thus, each leaf is associated with labels indicating all of the strings and starting positions of the associated suffix.

Generalized Suffix Trees

T_1 : xabxa

T_2 : babxba





Generalized Suffix Trees

How to build the suffix tree for a set $T = \{T_1, T_2, \dots, T_i\}$ in $O(m)$?

- (1) Build a suffix tree for $T_1\$$
- (2) Start from the root of the tree search for T_2 . Assume that i characters in T_2 are matched,
 - The suffix tree has implicitly encoded every suffix of $T_2[1..i]$
 - The suffix tree contains I_i for T_2
 - We can skip phase $1..i$ for T_2
- (3) Continue the Ukkonen's algorithm on T_2 in phase $i+1$
 - Walk up from the end of $T_2[1..i]$, ...
- (4) Until all T_i are included in the suffix tree.



4. Longest Common Substring (LCS) of Two Strings

Given strings S_1 and S_2 , find the LCS of them.

Different from *longest common subsequence problem*.

S_1 : *xabxa*

S_2 : *babxba*

LCS is *abx*




4. Longest Common Substring (LCS) of Two Strings

- Build a generalized suffix tree for S_1 and S_2
- If a leaf is from S_1 , then mark all its ancestors with 1.
- If a leaf is from S_2 , then mark all its ancestors with 2.
- The path-label of any node that is marked with both 1 and 2 is a common substring of S_1 and S_2 .
- Find the node that is labeled with 1 and 2, and has the greatest **string-depth** (number of characters on the path to it).

$T_1: xabxa$

T₂: babxba






4. Longest Common Substring (LCS) of Two Strings

- (1) $O(m)$ for building generalized suffix tree
- (2) $O(m)$ for calculating the string-depth of each node (e.g. Breadth first)
- (3) $O(m)$ for marking node with 1 or 2 (e.g. Depth first)
- (4) $O(m)$ finding the longest.



5. DNA Contamination Problem

DNA contamination: During laboratory processes, unwanted DNA inserted into the DNA of interest.



Contamination sources: Human, bacteria,...

DNA from Dinosaur bone: More similar to human DNA than to bird and crockodilian DNA



5. DNA Contamination Problem

S: DNA of interest

P: DNA of possible contamination source


If S and P share a common substring longer than ℓ , then S has been contaminated by P.

To find all common substrings of S and P that are longer than ℓ .

In general, P is set of DNA that are potential contamination sources.

$T_1: xabxa$

T₂: babxba





6. Common Substrings Of More Than Two Strings

Motivation

alignment

ALRDEATHDDF
SMTAEATHDSI
ECDQAATHEAS

regular
expression

A-T-H-[D,E]



6.Common Substrings Of More Than Two Strings

Problem statement: Given K strings whose lengths sum to n , let $l(i)$ be the length of the longest substring common to at least i strings, to compute a table of $K-1$ entries, where entry i give $l(i)$ and one of the common substrings of that length (and that is shared by at least i strings)

{sandollar, sandlot, handler, grand, pantry}

i	$l(i)$	substring
2	4	sand
3	3	and
4	3	and
5	2	an



6. Common Substrings Of More Than Two Strings

It can be solve in $O(n)$ time.

But, an easy algorithm that uses $O(kn)$ time first.

- Build a generalized suffix tree for the k strings giving each string a unique end marker.
- Each leaf belong to only one string
- For a node (v), let $c(v)$ be the number of distinct string identifiers that appear at the subtree below it.
- V is a vector with $V(i)$ denoting the length of the longest substring that occurs **exactly in i strings** (and a pointer to the node).
- From $V(i)$ compute $l(i)$,
- for $i=k; i>1; i--$
- if ($V(i)<V(i+1)$), then $l(i)= V(i+1)$
- else $l(i)= V(i)$


$$l$$
21



6. Common Substrings Of More Than Two Strings

Calculating $c(v)$ is the bottle neck.

Can't just count the number of leaves below it.


- For each node keep a C vector of k bits, with one bit correspond to one string.
- i^{th} is set to 1 if a leave that belongs to i^{th} string appear below the node
- The V vector of a parent is obtained by ORing the vectors of its children.
- n nodes.
- $O(Kn)$ in calculating $c(v)$.

Suffix Trees to DAGs

Space is a big
problem for suffix
trees.


S: xyxaxaxa\$

The subtree under p
is **isomorphic** to
that under q
except for leaf
labels



Suffix trees to DAGs



Directed acyclic graph (DAG)





Suffix Trees to DAGs

S: xyxaxaxa\$

P: xax



Suffix Trees to DAGs



If the subtrees under p and q are isomorphic (except leaf labels) and $\text{stringdepth}(p) > \text{stringdepth}(q)$, then

- Merge p into q , by adding a direct edge from $\text{parent}(p)$ to q
- Associate the directed edge with $d = \text{stringdepth}(q) - \text{stringdepth}(p)$
- When search for P in the S (text), let i be the leaf below the path labeled with P , if the directed edge is traversed then P occurs at $i+d$, otherwise P occurs at i .



Suffix Trees to DAGs

How to determine whether a subtree is isomorphic to another one?

Theorem 7.7.1

In suffix tree T the subtree below a node p is isomorphic to the subtree below a node q **if and only if**

- there is a directed **path** of suffix links from one node to the other node and
- the numbers of leaves in the two subtrees are equal.

A if and only if B

$B \rightarrow A$

$A \rightarrow B$


Ukkonen Algorithm

Suffix links

Let $x\alpha$ denote an arbitrary string, where x denotes a single character and α denotes a (possible empty) substring. For an **internal node** v with path-label $x\alpha$, if there is another **node** $s(v)$ with path-label α , then a pointer from v to $s(v)$ is called a **suffix link**, denoted as $(v, s(v))$.

The root has no suffix link from it.

If α is empty, then the suffix link points the root.





Suffix Trees to DAGs

$B \rightarrow A$

Only one suffix link

For every path from p to a leaf in its subtree, there is an identical path from q to a leaf in its subtree.




Suffix Trees to DAGs

$B \rightarrow A$

A **path** of suffix links

For every path from p to a leaf in its subtree, there is an identical path from q to a leaf in its subtree.




Suffix Trees to DAGs

$A \rightarrow B$

Either α is a proper suffix of γ
or γ is a proper suffix of α

There is a directed path of suffix
links from one node to the other.




Suffix Trees to DAGs

$A \rightarrow B$

Either α is a proper suffix of γ
or γ is a proper suffix of α

There is a directed path of suffix links from one node to the other.





Suffix Trees to DAGs

Let Q be the set of all pairs (p,q) such that there is a suffix link from p to q .

While there is a pair (p,q) in Q

 Merge p into q ;

 Remove (p,q) ;

The merge of the pairs can be done in arbitrary order.

In practice, we can start merge in a top-down approach (depth-first).



Suffix Arrays—more space reduction

Given a m -character string T , a suffix array for, called Pos , is an array of integers in the range 1 to m , specifying the lexicographic order of the m suffixes of string T .

$Pos[i]$ lexically less than $Pos[i+1]$

mississippi

pos 11,8,5,2,1,10,9,7,4,6,3



Suffix tree to suffix array

- In $O(m)$ time
- Lexical depth-first search



Pattern searching using suffix arrays

Observation: If p occurs in T then all the locations of those occurrences will be grouped consecutively in Pos .

$P=issi$

$T=mississippi$



Pattern searching using suffix arrays

Basic idea: Binary search

$O(n \log m)$ (worst)

$O(n + \log m)$ (expected)



A simple accelerant

L and R are left and right boundaries of the "current search interval".

Query will be made at $M = (L+R)/2$ of Pos.

l: the length of the longest prefix of Pos(L) that match a prefix of P

r: the length of the longest prefix of Pos(R) that match a prefix of P

$l_{mr} = \min\{l, r\}$

Compare P and Pos(M) starting from position $l_{mr} + 1$ of the two string.

$O(n \log m)$



A super accelerant

- Lcp (i,j): length of the longest prefix of Pos(i) and Pos(j)
- Use Lcp(L,M), Lcp(M,R)
 - Suppose $l > r$,
 - If $\text{Lcp}(L,M) > l$, $L \leftarrow M$, and l, r unchanged
 - If $\text{Lcp}(L,M) < l$, $R \leftarrow M$, $r = \text{Lcp}(L,M)$
 - If $\text{Lcp}(L,M) = l$, comparison of P and Pos(M) starting at $l+1$.
- $O(n + \log m)$



To obtain Lcp (i,j)

- Lcp (i,i+1) for i=1 to m-1
 - Lexical depth first search
- For any $i < j$, Lcp(i,j) is the smallest value of Lcp(k,k+1), where, k=i to j-1
 - Lexical depth first search in a complete binary tree